

OSSAMA TAHA

BI & DATA ANALYTICS ENGINEER

📍 6th of October City, Egypt | ✉ ossama.y.taha@gmail.com | 📞 +201026458878 | 🌐 linkedin.com/in/ossamataha
🔗 github.com/OssamaTaha

SUMMARY

BI & Data Analytics professional with a B.Com in Accounting and 2+ years of experience building production-grade data pipelines, automated reporting systems, and executive dashboards. Proven track record of reducing reporting cycles by 10x through end-to-end automation, enabling data-driven strategic decisions for sales and operations teams. Combines deep business domain knowledge with strong technical skills in Python, SQL, PostgreSQL, Apache Airflow, dbt, and cloud infrastructure.

EXPERIENCE

Nile for Cables and Electronics , BI & Sales Data Analyst	Giza, Egypt
- Reduced reporting cycle from days to hours by designing and deploying a semi-realtime automated reporting system, directly improving sales team responsiveness and strategic decision-making - Led quarterly sales performance evaluations using consolidated KPIs, providing actionable insights that drove revenue-focused business decisions - Built executive dashboards supporting sales forecasting and real-time KPI monitoring, enabling stakeholders to track performance without manual data pulls - Designed end-to-end ETL pipelines using PostgreSQL and Dolt (version-controlled SQL database), ensuring data integrity and audit trail for critical financial metrics - Developed a local AI chatbot using open-source LLMs for natural language querying of business data, empowering non-technical stakeholders to self-serve analytics	Nov 2024 – present 1 year 7 months
Samsung Innovation Campus , Data Engineer	Giza, Egypt
- Developed scalable data processing workflows using Apache Spark and Kafka, supporting business intelligence and analytical use cases - Performed exploratory data analysis and feature engineering to enable data-driven business insights - Deployed production data workflows on AWS cloud infrastructure (EC2, S3), ensuring reliable and cost-efficient operations	July 2024 – Oct 2024 4 months

PROJECTS

dbt Data Warehouse – Star Schema & Analytics Layer Problem: Raw operational data lacked structure for analytics, forcing manual data wrangling for every report Process & Solution: Designed star schema with 15+ dbt models across staging/marts layers, SCD Type 2 tracking, incremental processing, and 99.9% data quality via automated testing Business Impact: Reduced query complexity by 70%, eliminated manual data prep, enabled self-service analytics
Atlas Analytics Platform Problem: Siloed data across departments required manual updates with no unified visibility Process & Solution: Built full-stack platform with Python backend, automated ETL pipelines, role-based dashboards, deployed on AWS with CI/CD Business Impact: Enabled data-driven decisions across 3 departments, reduced manual reporting by 80%
Airflow ETL Pipeline – Multi-Source Data Orchestration Problem: Manual API data extraction was error-prone with no scheduling or failure recovery Process & Solution: Designed 4 production DAGs with task dependencies, dynamic mapping, retry logic, extracting from 3 API sources into PostgreSQL with real-time monitoring Business Impact: Fully automated daily/weekly ingestion with zero manual intervention
Agentic AI Workflow – Multi-Agent DataOps Automation Problem: Data operations required constant manual monitoring and debugging across multiple systems Process & Solution: Built multi-agent system with orchestrator pattern, 5 specialized agents, tool integration, and persistent memory using LangChain/LangGraph Business Impact: Cut pipeline failure resolution time by 70%, automated 90% of routine monitoring
AWS Cloud Data Pipeline – Serverless ETL with IaC Problem: On-premise infrastructure lacked scalability and cost efficiency for growing data Process & Solution: S3 data lake + Glue ETL + Athena queries + Step Functions orchestration, fully deployable via Terraform IaC Business Impact: 90% cost reduction vs traditional infra, sub-second query performance, zero server management
Kafka Streaming Pipeline – Real-Time Data Processing Problem: Batch processing couldn't meet real-time analytics needs for stock monitoring and telemetry Process & Solution: Event-driven architecture with Kafka producers, consumer groups, windowed aggregation, dead letter queues, live dashboard, and PostgreSQL sink Business Impact: Sub-second data latency, automated threshold alerting, real-time operational visibility

SKILLS

- BI & Data Visualization:** Power BI, DAX, Tableau, Excel, Streamlit, Plotly, Dash
- Data Pipeline & Architecture:** ETL/ELT Pipelines, Apache Airflow, dbt, Apache Spark, Kafka, Data Modeling (Star Schema), PostgreSQL, MySQL
- Cloud & Infrastructure:** AWS (S3, Glue, Athena, Lambda, Step Functions, EC2), Terraform, Docker, Docker Compose, Linux Server Administration, Nginx, GitHub Actions (CI/CD)
- Programming & Scripting:** Python, SQL, Bash, JavaScript, TypeScript, HTML/CSS
- AI & Machine Learning:** LangChain, LangGraph, Hugging Face Transformers, Scikit-learn, TensorFlow, NLP, Local LLM Deployment
- Data Quality & Monitoring:** Great Expectations, Data Validation, Schema Testing, Anomaly Detection, Pipeline Monitoring
- Automation & Tools:** Git, n8n, OpenClaw, KiloCode, Hermes, Dolt, Jupyter, VS Code

EDUCATION

- Al-Azhar University**, Accounting & Business Management Sept 2019 – June 2023
- Bachelor of Commerce**
- Self-taught in programming, data engineering, and cloud technologies through professional certifications and hands-on projects

CERTIFICATIONS

- Google Data Analytics Professional Certificate**
- Data cleaning, analysis, visualization with spreadsheets, SQL, Tableau, and R
- IBM Data Engineer Professional Certificate**
- ETL pipelines, data warehousing, big data (Spark, Hadoop, Kafka), cloud deployment
- IBM Data Science Professional Certificate**
- Machine learning, data analysis, Python, SQL, applied data science capstone projects
- Samsung Innovation Campus Big Data Engineering Certificate**
- "Top Achiever (#1) in the Big Data Engineering track"
 - Intensive training in big data infrastructure, distributed computing, and data engineering
- AWS Certified Data Engineer – Associate**
- S3, Glue, Athena, Lambda, Step Functions, and Terraform (In Progress)
- iMBA Business Analysis Specialization (In Progress)**
- Coursera iMBA from UIUC — business process modeling, requirements engineering, data-driven decisions